

Belief propagation

Data Intensive Programming

Andrea Scotti, Sandrine Idlas

KTH

October 28, 2019

Belief propagation

Message passing algorithm

- Used on probabilistic graphical models.
- Two types of nodes:
 - ▶ Variable nodes represent a stochastic variable present in the model.
 - ▶ Factor nodes contain a factored representation of the joint probability of its adjacent nodes.
- Algorithm aim: evaluate the marginal probability at every node.

Project aim

- Implement this algorithm in a distributed way.
- Use the model on trees.
- Stochastic variable is binary, with values 0 and 1.

Belief propagation


Joint distribution

$$P(A, B) = P(A|B)P(B):$$



Alternative representation


- circle nodes = variables
- dark square = factor node



- Marginalization:
 $P(B) = \sum_A P(B|A)$
- Message from A to the factor node: $\mu_{A \rightarrow AB} = P(A)$.
- Message from the factor node to B : $\mu_{AB \rightarrow B} = \sum_A \psi_{AB} \cdot \mu_{A \rightarrow AB}$.


Belief propagation

Chain



- The joint distribution is $p(a, b, c, d) = f_1(a, b)f_2(b, c)f_3(c, d)f_4(d)$.
- $\mu_{d \rightarrow c}(c) = \sum_d f_3(c, d)f_4(d)$.
- $\mu_{c \rightarrow b}(b) = \sum_c f_2(b, c)\mu_{d \rightarrow c}(c)$.
- $p(a) = \sum_b f_1(a, b)\mu_{c \rightarrow b}(b)$

Branching



- Model:
 $p(a|b)p(b|c, d)p(c)p(d)p(e|d) = f_1(a, b)f_2(b, c, d)f_3(c)f_4(d, e)f_5(d)$.
- Marginalization: $\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d)f_3(c)f_5(d) \sum_e f_4(d, e)$.

Belief propagation

Algorithm for updating messages as a function of incoming messages:

① Initialisation:

- ▶ The messages from leaf node factors are initialized to the factor.
- ▶ The messages from leaf variable factors are initialized to unity.

② Factor to variable:

- ▶ Formed by summing the product of incoming node-to-factor messages.
- ▶ $\mu_{x \rightarrow b}(x) = \prod_g \mu_{g \rightarrow x}(x)$, node x is excluded from the product.

③ Variable to factor:

- ▶ Product of incoming factor-to-node messages, does not imply a summation.
- ▶ $\mu_{f \rightarrow x}(x) = \sum_{X_f} \psi(X_f) \prod_y \mu_{y \rightarrow f}(y)$, summation over all states of all variables, except the current one.

Graphs in GraphX

- Graphs represented as VertexRDD and EdgeRDD.
- Executed using Pregel.
- Supersteps: sequences of iterations for message passing between distributed vertices.
- Allows to store trees bigger than the memory of a single machine.
- Can compute big trees since computation is shared between machines.

Input

- Fetched from a text file.
- Variable node ID: from 0 to the number of variables minus one.
- Factor node ID: shifted by the number of variable nodes.
- Each variable node x can have two different realizations (0 or 1), with probabilities $p(x = 0)$ and $p(x = 1)$.

Example of factor node with two neighbours:

```
0 1
0.3 0.7 0.7 0.3
```

- Adjacent variable nodes ID: 0 and 1.
- Distribution:
$$\begin{bmatrix} 0.3 & 0.7 \\ 0.7 & 0.3 \end{bmatrix} = \begin{bmatrix} p(0,0) & p(0,1) \\ p(1,0) & p(1,1) \end{bmatrix}$$

Supersteps in Pregel

Running Pregel

```
val finalGraph = graphWithLeaves.pregel  
    (initialMsg, maxIterations, EdgeDirection.Out)  
    (vprog, sendMsg, mergeMsg)
```

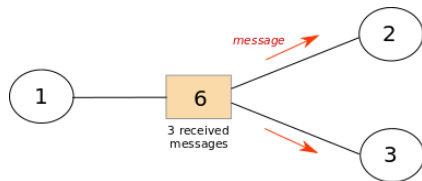
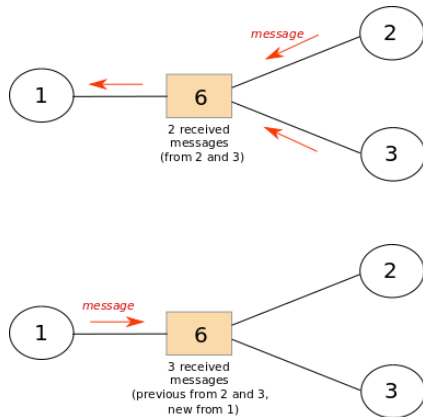
Functions mergeMsg, vprog and sendMsg

- ❶ Merge (function mergeMsg): concatenates two lists (incoming messages to a factor node).
- ❷ Apply (function vprog):
 - ▶ Adds new messages to the list of collected messages of the node.
 - ▶ Stores the source of a message in certain cases.
- ❸ Scatter (function sendMsg):
 - ▶ Called for each edge of every scattered node.
 - ▶ Leaf nodes: only send messages at the beginning of the algorithm.
 - ▶ Internal nodes: send messages once they have received enough messages.
 - ▶ Computes messages to be sent using the message list created in vprog.

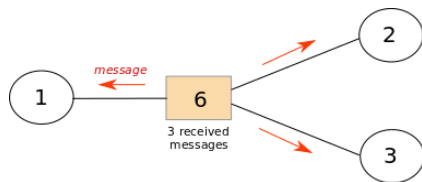

Initialization

- **Variable leaf nodes:** initialized to $(1, 1)$, or $(0, 0)$ for log.
- **Factor node:** initialized to the factor itself.

Choosing node destination



Choosing node destination



Output

After the algorithm has ended:

- Filter results by node type, keep only variable nodes.
- Compute probabilities for realizations 0 or 1.
- For each node, print the result:

(Node ID, ($p(0)$, $p(1)$, Normalization Constant))

Normalization constant

- Used to double-check that the algorithm functioned properly.
- Should be the same for every node \implies information was passed properly through the network.

Results when using log

- e^x , where $x = \log(p(x))$, can be too small \implies retrieving the normalization constant is not possible.
- Print a list of non-normalized log probabilities:

(Node ID, ($\log(p(0))$, $\log(p(1))$))

Method

Implementation

- Python script to understand the algorithm.
- Inputs generated by Python script.
- Implement the algorithm on GraphX.
- Run it in a distributed way using the cluster Hopsworld.

Debugging issues

- Cannot easily print with GraphX.
- Workaround: Accumulator function to accumulate the states during computation.

Next step

- Split the input file in several input files in Hadoop.
- Potential problems:
 - ▶ Need to create a local array that could contain all the nodes during the reading of the input.
 - ▶ Risk of runtime error during execution.
- Instead: use `spark.kryoserializer.buffer.max = 1024M` \Rightarrow could run half a million of nodes.

Dataset

- synthetic dataset
- python script called *factorTreeGenerator.py*
- generates a random tree rooted given:
 - ▶ total number of nodes n
 - ▶ minimum and maximum number of children for each node
- rooted in the node with ID 0

Results

- Number of nodes 100 \Rightarrow distributed algorithm with GraphX slower than centralized Python script.
- Computation time: several minutes, even in a big cluster. Example:
 - ▶ tree with 500000 nodes
 - ▶ number of children for each node uniformly between 1 and 10
 - ▶ cluster configurations:

Hours to shutdown ⓘ 6

Driver memory (MB) ⓘ 16384

Driver virtual cores ⓘ 8

Executor memory (MB) ⓘ 16384

Executor virtual cores ⓘ 8

Set min and max number of executors with the Slider below.

0 64 128

Initial executors ⓘ 68

- ▶ **Computation time:** 20 minutes.

Results

- Trees of similar size: high influence of topology on the running time performance of the algorithm:
 - ▶ On a simple chain with 1 child per parent node, the algorithm runs very slowly in a distributed way.
 - ▶ Distributed computation more performant the more children there are per node, but the size of the distribution array exponentially increases with the number of neighbors k as 2^k .